

Viyog: Separating Adversarial and Out-of-Distribution

Aman Singh, Abhinav Kumar*, Rishab Kashyap[†], Soyeong Park[‡], Ameya Gurjar[‡], Yohan Ko, Hwisoo So*,
Uiwon Hwang, Kyoungwoo Lee, and Aviral Shrivastava

Abstract—Safety-critical systems must respond differently to two anomalies: out-of-distribution (OOD) inputs, which call for abstention, and adversarial (ADV) attacks, which demand rejection. Existing detectors collapse both into a single anomaly score. We present Viyog, a training-free plug-in second stage that separates OOD from ADV by reading the dormant-band roughness of the first convolutional activation—how spatially jagged its quietest channels are (a magnitude-normalised total variation, i.e. average change between neighbouring pixels), computable from the forward pass already in progress and storing only ≈ 0.3 KB ($O(C)$) versus 4.5–40 MB for feature-distance baselines.

Gradient attacks inject broadband high-frequency residue into the quietest first-layer channels; natural inputs (ID and OOD alike) leave them smooth. Viyog exploits this asymmetry: it detects adversarials at AUROC 0.966 (where 1.0 is perfect and 0.5 is chance) while logit detectors are blind (≤ 0.41); combined with logit scores it gives the first deployable three-way ID/OOD/ADV separator (balanced accuracy 0.798 over 20 architectures; 0.863 on GTSRB). We prove scale-invariance (P1)—so norm-preserving attacks achieve 0% evasion—and a concrete separation bound (P2/P3). The detector generalises across 3 datasets, 20 architectures, 4 attacks, and 1D non-vision signals, adds no parameters, and runs in a single forward pass.

Index Terms—Out-of-distribution detection, adversarial robustness, feature norms, deep neural networks, safety-critical systems, embedded AI, architecture generalization.

I. Introduction

Deep neural networks (DNNs) achieve state-of-the-art performance across safety-critical domains including autonomous driving, medical imaging, and embedded vision systems. Deployed systems must therefore detect and respond appropriately to abnormal inputs before acting on model predictions—a requirement that standard classification pipelines do not satisfy [1].

Two kinds of anomalous input require opposite responses (Fig. 1): OOD inputs fall outside the training distribution and call for abstention, while adversarial (ADV) inputs are crafted to force misclassification and demand active rejection. Conflating them wastes defences




Fig. 1: The core problem: OOD and ADV inputs require opposite responses, yet both fool the classifier. (1) Input space: an OOD (out-of-distribution) input lies outside the training cluster; an adversarial (ADV) input is deliberately crafted to sit inside the cluster while forcing a wrong prediction. (2) Classifier: both produce anomalous logits that existing output-side detectors cannot distinguish. (3) Viyog: reads the dormant-band roughness $V(x)$ of the first convolutional activation map—a gradient-free, 0.3 KB statistic—and separates the two. (4) Response: OOD $\rightarrow$ abstain / flag for review; ADV $\rightarrow$ reject / security alert. Conflating them wastes adversarial defences on benign inputs or silently accepts attacks.

or silently accepts attacks. No existing lightweight detector makes this distinction; we show the signal is hidden in the first layer.

We identify a consistent, architecturally-stable asymmetry in first-layer activation statistics that enables this separation at no additional cost. Gradient-based adversarial attacks inject a broadband residue—noise spread across all spatial frequencies—into the first convolutional or projection layer: the attack spends its norm budget on a high-curvature surface, so the per-pixel perturbation passes through even the quietest (“dormant”) filters as high-frequency structure. Natural inputs (ID and OOD), unconstrained by any such objective, leave those low-gain channels smooth [2], [3]. An earlier, weaker form of Viyog exploited this asymmetry with the raw first-layer L_∞ norm (the L_∞ -form): adversarials fall below the ID mean and OOD at or above it. But the L_∞ -form is regime-bounded—strong only for far-OOD (inputs far from the training distribution, e.g. a different dataset) and near-chance for near-OOD (visually similar shifts) under a strict cross-model protocol. The deployable signal is instead the relative roughness of the dormant channels: a magnitude-normalised total-variation (TV) statistic (Viyog, Sec. IV) that measures shape (spatial roughness), not magnitude. Being a 0-homogeneous ratio—scale-invariant, so rescaling all activations leaves it unchanged—it is provably immune to norm-preserving attacks and recovers the near-OOD regime the L_∞ -form loses. It is complementary to logit

A. Singh, A. Kumar, R. Kashyap, A. Gurjar, and A. Shrivastava are with Arizona State University, Tempe, AZ, USA (e-mail: asing651@asu.edu; akuma294@asu.edu; rkashya6@asu.edu; agurjar2@asu.edu; aviral.shrivastava@asu.edu).

S. Park, Y. Ko, and K. Lee are with Yonsei University, Seoul, Republic of Korea (e-mail: psy980823@yonsei.ac.kr; Yohan.Ko@yonsei.ac.kr; kyoungwoo.lee@yonsei.ac.kr).

H. So is with Kyungpook National University, Daegu, Republic of Korea (e-mail: hwisoo.so@knu.ac.kr).

U. Hwang is with Ewha Womans University, Seoul, Republic of Korea (e-mail: uiwon.hwang@ewha.ac.kr).

*Corresponding authors: Abhinav Kumar (e-mail: akuma294@asu.edu) and Hwisoo So (e-mail: hwisoo.so@knu.ac.kr).

[†]These authors contributed equally to this work.

OOD scores and yields three-way ID/OOD/ADV separation (Sec. VI); the phenomenon persists across CNN and transformer architectures and across 32×32 and 224×224 inputs. We ground this in a formal Proposition (Section III) that proves both halves: the dormant-band roughness inflates under any broadband attack, yet stays invariant to activation magnitude.

This paper makes five core contributions:

C1 — The Viyog scoring statistic. A gradient-free first-layer statistic—the magnitude-normalised total variation of the quietest (dormant) channels—that separates ADV from OOD where the raw L_∞ norm cannot (mechanism confirmed on real ResNet-50 features, Fig. 3). The continuous statistic $V(x)$ is itself the graded degree-of-anomaly readout; the monotone routing squash is retained only to shape the binary operating curve, not as the severity scale. The raw L_∞ score is kept as the motivating baseline.

C2 — Theory for the shape statistic. A formal Proposition (Sec. III-B) proving that the dormant-band roughness statistic is scale-invariant—so norm-preserving attacks cannot move it—and that it inflates under the exact affine residue $e_c = w_c * \delta$ any broadband attack injects, for any affine first layer (convolution or ViT/Swin patch embedding) and with no activation-curvature assumption (a 1-Lipschitz readout suffices). The ID/OOD/ADV ordering is stated as a conditional consequence and an experimentally validated prediction, not a worst-case guarantee; a general-affine operator-norm envelope (Appendix A) covers the raw- L_∞ baseline.

C3 — Regime characterisation and a complementary dormant-band detector. Delimiting where the global L_∞ mechanism holds and recovering the lost near-OOD/strong-attack regime with Viyog, we obtain the first deployable three-way ID/OOD/ADV separation: across 20 architectures and two datasets Viyog is complementary to logit OOD scores, and a lightweight supervised head over the full detector panel reaches balanced accuracy 0.80 (CIFAR-100, 20 archs) / 0.86 (GTSRB, 6 archs), the logit+Viyog pair alone 0.76/0.77, at ≈ 0.3 KB of added state—validated on ConvNeXtV2, ViT-B/16, and Swin-T at 224×224 .

C4 — Comprehensive empirical validation. Experiments across five datasets and four L_∞ attacks. Under the original far-OOD single-direction protocol the raw L_∞ score outperforms ten established baselines by up to +24.9 AU-ROC points; the stricter directionless re-evaluation—which disqualifies any score whose high/low meaning flips between models—over 20 architectures is reported in Table III.

C5 — Adaptive attack analysis. Norm-preserving evasion achieves 0% evasion; a both-aware adaptive attacker pays an 8% success cost; the stochastic Expectation-over-Transformation (EOT) defence—randomising which dormant channels are read so the attacker cannot lock onto a fixed target—restores the panel to 0.82 (Fig. 14).

TABLE I: Prior anomaly detectors and their OOD–ADV separation capability. All existing families fail at second-stage OOD-vs-ADV separation. † Viyog is the only method that scores both T2 and T3 above chance while remaining gradient-free and state-minimal.

Family	Examples	Score	Sep. OOD/ADV?	Cost
Confidence	MSP, Energy, ODIN	logit	✗ both high conf.	$O(1)$
Distance	Maha, KNN, ViM	feat. dist.	✗ ADV inside manifold	$O(D^2)$ – $O(ND)$
Uncertainty	MCD	post. var.	✗ single score	30× fwd
Joint	UNICAD [7]	multi	partial	grad. at inf.
Viyog†	this work	shape V	✓ T2 0.97 / T3 0.82	0.3 KB

II. Background and Related Work

A. OOD Detection Methods

Table I categorises prior work; surveys are in [4], [5]. All methods collapse OOD and ADV into a single anomaly score; none provides a deployable second-stage separator. Szyk et al. [6] show that rankings are protocol-sensitive, motivating our multi-seed evaluation (Sec. V).

B. Adversarial Attacks, Defences, and Early-Layer Effects

Having surveyed OOD scoring, we turn to the second anomaly class Viyog must separate from it. Adversarial attacks [8] solve $\min_\delta \|\delta\|_p$ s.t. $f(x+\delta, \theta) \neq y$. Common attacks include FGSM [9], PGD [10] (both L_∞), and L_2 -constrained CW [11] and DeepFool [12]. A fundamental tension exists between robustness and accuracy [13], [14].

Defences against adversarial attacks fall into several families: input preprocessing [15], feature denoising [16], generative purification [17]–[20], and adversarial training [10], [14]. Ensemble and Jacobian-based methods provide additional robustness but at increased inference cost. None of these defences address the OOD/ADV disambiguation problem; they assume the threat class is known in advance, precisely the assumption Viyog is designed to relax.

Early-layer effects are well documented: the first convolutional layer is disproportionately critical for robustness [3] and acts as an implicit noise filter [2] (the front-neuron concept formalised here); pruning low-activation neurons improves robustness by forcing reliance on high-magnitude responses [21], and restoring suppressed activations [22] presupposes the suppression Viyog detects. Filter-statistic [23] and frequency-domain [24], [25] analyses likewise distinguish adversarial perturbations from natural shift at early layers, but none exploits the OOD/ADV asymmetry. We sharpen these into a formal Proposition (Section III): adversarial perturbations systematically suppress the highest-magnitude activations in the earliest convolutional layer, the geometric foundation of Viyog.

C. Joint OOD and Adversarial Detection

A small body of work attempts the OOD–ADV separation directly, but at a cost Viyog avoids. Zhang et al. [7] address OOD–ADV separation via adversarial gradient attribution, requiring gradient computation at inference time. Feng et al. [26] propose a unified open-world detection perspective but require joint training with




Fig. 2: The Viyog two-stage deployment pipeline. A normal forward pass drives a first-stage OOD gate (logit-based scores such as Energy [29] or MSP (Maximum Softmax Probability [30])) and, via a zero-cost hook, exposes the first-conv activation map $a = f_0(x)$. Flagged inputs enter Viyog, which computes $V(x)$ —the dormant-band roughness (total-variation shape statistic) of a —and routes to ID (pass), OOD (abstain / flag), or ADV (reject / alert), adding only 0.3KB of state and no extra forward pass.

anomaly labels. Walkowiak et al. [27] combine outlier-ness scores and feature extraction, while UNICAD [28] unifies attack detection and novel-class identification in a single framework. Detection methods based on feature statistics [23] and disentangled representations have also been proposed, but none provides a single gradient-free scalar operable as a plug-in second stage. Unlike gradient-attribution or joint-training approaches, Viyog operates in a single gradient-free forward pass and serves as a plug-in second stage atop any existing non-ID detector, without architectural modification. The Viyog statistic and its unsupervised OOD-vs-ADV threshold need no labelled anomaly data; the optional three-way head adds a lightweight supervised calibration over a small labelled set (Sec. VI). This efficiency advantage is quantified in the system analysis that follows.

D. System Efficiency and Deployment Motivation

For embedded systems inference cost is a primary constraint [1], yet prior detectors are expensive (ODIN and gradient-norm backward passes, a full Mahalanobis covariance, MCD’s 30 forward passes ($\sim 30\times$ a single inference), diffusion denoising). Viyog adds zero parameters, no gradient computation, and only a first-layer activation hook (read as a dormant-band shape statistic); the pipeline below grounds this in deployment.

In the Viyog two-stage pipeline (Fig. 2), a standard inference pass is followed by a first-stage non-ID detector that flags inputs outside the in-distribution region, and Viyog partitions them into OOD and ADV using only the first-layer activation hook already computed in the forward pass (read as the dormant-band roughness statistic V). Without this second stage, a system cannot distinguish a novel but benign environment from an active adversarial attack. Prior joint approaches [17], [18] cannot make this distinction at deployable cost; the downstream routing consequences are detailed in Sec. IX.

III. Theoretical Motivation and Analysis

We formalise why gradient attacks leave a detectable signature at the first layer, and prove the properties of the dormant-band roughness statistic $V(x)$ (Prop. 1) that make it the deployed detector: inflation under any broadband attack, invariance to activation magnitude, and a concrete separation guarantee under three testable conditions validated by the empirical results of Section VI. The argument is rooted in the convolutional nature of the first layer shared by all architectures evaluated: strided convolutions in ResNet [31], DenseNet [32], and ConvNeXtV2 [33]; and patch-embedding projections (which are mathematically equivalent to strided convolutions) in ViT [34] and Swin-T [35]. The insight that high-frequency and norm-altering perturbations are concentrated at early layers is supported by prior frequency-domain analyses [15], [24], [25].

A. Why the First Layer?

Natural images are smooth at the first layer (Premise A1). Natural image power spectra decay as $P(f) \propto f^{-2}$: low-frequency components (slow gradients, textures) dominate; rapid pixel-to-pixel oscillations are negligible. Each first-layer channel c produces a feature map $a_c \in \mathbb{R}^{H \times W}$ —the response of filter w_c to input x . Because natural images are spectrally smooth, a_c is smooth too: its Total Variation $\text{TV}(a_c)$ (mean pixel-to-pixel jump, defined formally in Eq. (1)) is small relative to the channel’s mean magnitude $|a_c|$. Crucially, this holds for all natural inputs—both ID and OOD—since both are photographs of the physical world.

Adversarial attacks inject broadband high-frequency noise (Premise A2). Gradient attacks add a perturbation δ to the input to maximally raise the loss. The loss gradient $\nabla_x \mathcal{L}$ carries energy at all spatial frequencies; FGSM takes its elementwise sign to produce a $\pm\epsilon$ broadband checkerboard pattern; PGD accumulates such steps. The L_∞ budget $\|\delta\|_\infty \leq \epsilon$ imposes no spectral shaping, so every frequency can independently saturate its ϵ slot. At the first layer, each filter w_c convolves with δ to produce a per-channel residue $e_c = w_c * \delta$ —an additive high-frequency noise map. Because δ is spectrally broadband, e_c has high Total Variation regardless of which specific attack generated δ .

Why a normalised ratio, not a raw norm (0-homogeneity). A raw L_∞ norm measures magnitude, not shape. OOD images can produce larger or smaller first-layer activations than ID images, so $\|a_c\|_\infty$ is direction-inconsistent; an attacker can further append a norm-preserving penalty to keep $\|a'_c\|_\infty \approx \|a_c\|_\infty$ while misclassifying. The normalised ratio $\text{TV}(a_c) = \text{TV}(a_c) / (|a_c| + \epsilon)$ (small $\epsilon > 0$ for stability) fixes both problems: TV and $|\cdot|$ are both 1-homogeneous, so their ratio is 0-homogeneous—scale-invariant (P1 below). An attacker who rescales activations cannot move TV; the only escape is to smooth δ , conflicting with misclassification under a tight ϵ -budget.




Fig. 3: The Viyog mechanism confirmed on real ResNet-50 features (CIFAR-100, $n=2000$ samples per class). (a) Dormant-band channels (bottom 10% alive, green) have the smallest denominators, amplifying injected roughness. (b) ID and OOD co-cluster at low $V(x)$; ADV-PGD spikes high (AUROC shown; dashed line = calibrated threshold). (c) Scale-invariance (P1): doubling activation magnitude leaves V unchanged—only the adversarial roughness ratio is elevated.

B. The Dormant-Band Roughness Statistic

The statistic. Let $a = f_0(x) = W*x + b \in \mathbb{R}^{C \times H \times W}$ be the first-layer activation map (C channels, spatial size $H \times W$), with $a_c \in \mathbb{R}^{H \times W}$ the feature map of channel c , filter w_c , and bias b_c . Define:

$$V(x) = \frac{1}{|\mathcal{B}|} \sum_{c \in \mathcal{B}} \widetilde{\text{TV}}(a_c), \quad \widetilde{\text{TV}}(a_c) = \frac{\text{TV}(a_c)}{|a_c| + \epsilon}, \quad (1)$$

$$\text{TV}(a_c) = \frac{1}{2} (|\Delta_h a_c| + |\Delta_w a_c|),$$

where $\Delta_h a_c[i, j] = a_c[i+1, j] - a_c[i, j]$ and $\Delta_w a_c[i, j] = a_c[i, j+1] - a_c[i, j]$ are pixel differences, $\bar{u} = \frac{1}{HW} \sum_{i,j} u_{ij}$ is the spatial mean, and $\epsilon = 10^{-6}$ prevents division by zero. The dormant band is $\mathcal{B} = \text{bottom-}p\% \text{ of alive channels}$, where $\mathcal{A} = \{c : \mathbb{E}_{\text{ID}}[a_c] > \kappa\}$ ($\kappa = 10^{-4}$) excludes permanently silent channels.

Proposition 1 below formalises three properties of V . The dormant band amplifies each property: its small denominator $|a_c|$ magnifies any roughness numerator increase while leaving natural inputs bounded. Fig. 3 confirms this on real ResNet-50 features: $V(x)$ separates by type of signal (natural vs. attack), not by size of activation.

Proposition 1 (Relative-roughness inflation of a quiet affine channel). Fix a dormant alive channel $c \in \mathcal{B}$ with filter w_c . Write $a_c = w_c * x + b_c$ (clean feature map) with mean magnitude $\nu_c := |a_c| > 0$. For an adversarial input $x' = x + \delta$ (x_{adv} ; natural inputs written x_{nat}), the per-channel residue $e_c = w_c * \delta$ is exact by affinity (no Taylor or sign assumption). Write $a'_c = a_c + e_c$, $r_c := |e_c|/\nu_c$ (relative residue size), and $\rho_c := \text{TV}(e_c)/|e_c|$ (residue roughness). In plain words: ν_c is how loud the channel normally is, r_c how large the attack's leftover is relative to that, ρ_c how jagged that leftover is, and β_c (below) the small natural-image jaggedness. Assume three testable empirical premises: leftmargin=2em

- (A1) Natural smoothness: $\text{TV}(a_c) \leq \beta_c \nu_c$, $\beta_c \ll 1$, for all natural inputs (ID and OOD)—the $1/f^2$ spectral roll-off (Sec. III-A).
- (A2) Broadband residue: $\rho_c \geq \rho_{\min} > 0$ on adversarial inputs—high-frequency injection shared by FGSM/PGD/CW/DeepFool, decoupled from sign $\nabla_x \mathcal{L}$.
- (A3) Non-negligible residue: $r_c \geq r_{\min} > 0$ on adversarial inputs—the residue is detectable relative to the dormant channel's quiet baseline.

Then: leftmargin=2em

- (P1) Scale-invariance (proven, exact): for any $\lambda_c > 0$, $\widetilde{\text{TV}}(\lambda_c a_c) = \widetilde{\text{TV}}(a_c)$ at $\epsilon \rightarrow 0$. V measures shape, not magnitude: rescaling activations via any norm penalty leaves V unchanged.
- (P2) Two-sided bound (proven, attack-agnostic):

$$\frac{\rho_{\min} r_{\min}}{1 + r_c} - \beta_c \leq \widetilde{\text{TV}}(a'_c) \leq \beta_c + \rho_c r_c + O(\epsilon/\nu_c).$$

The residue satisfies the universal envelope $\text{TV}(e_c) \leq \|w_c\|_1 \|\nabla \delta\|_1$ ($\|\nabla \delta\|_1$: spatial mean of δ 's gradient magnitude; empirically $\sim 2 \times$ slack).

- (P3) ADV>OOD separation (conditional consequence): if $\rho_{\min} r_{\min} > 2\beta_c(1 + r_c)$ on $\mathcal{B}$, then $V(x_{\text{adv}}) > V(x_{\text{nat}})$ with per-channel margin $\frac{\rho_{\min} r_{\min}}{1 + r_c} - 2\beta_c \geq 0$, increasing in r_c .

P1 and the P2 inequalities are proven; P3 is a conditional consequence of A1–A3 and is experimentally validated (T2 AUROC ≈ 0.97 , T3 ≈ 0.82 , 20 architectures), not a worst-case guarantee. Post-activation readouts (ReLU/GELU, 1-Lipschitz) preserve P2; their curvature enters only the magnitude denominator, to which P1 makes V insensitive.

Proof sketch (full proof in Appendix A). P1. TV and $|\cdot|$ are both 1-homogeneous, so $\widetilde{\text{TV}}(\lambda a_c) = \frac{\lambda \text{TV}(a_c)}{\lambda |a_c| + \epsilon} \rightarrow \widetilde{\text{TV}}(a_c)$ as $\epsilon \rightarrow 0$. Consequence: the only escape for an attacker is to make δ smooth, conflicting with maximising the loss.

P2. TV is subadditive and satisfies a reverse triangle inequality:

$$\begin{aligned} \text{TV}(a_c + e_c) &\leq \text{TV}(a_c) + \text{TV}(e_c), \\ \text{TV}(a_c + e_c) &\geq \text{TV}(e_c) - \text{TV}(a_c). \end{aligned}$$

Substitute A1 ($\text{TV}(a_c) \leq \beta_c \nu_c$) and A2 ($\text{TV}(e_c) \geq \rho_{\min} r_c \nu_c$), divide by $|a'_c| \leq \nu_c(1 + r_c)$ (from A3).

P3. P2's lower bound minus A1's upper bound gives gap $\frac{\rho_{\min} r_{\min}}{1 + r_c} - 2\beta_c > 0$. Averaging over $\mathcal{B}$ gives $V(x_{\text{adv}}) > V(x_{\text{nat}})$. OOD stays low because A1 holds for all natural inputs and P1 prevents magnitude shifts from raising V . $\square$

In plain words. Natural images (ID and OOD alike) leave first-layer dormant channels smooth and quiet— V is low. Adversarial δ injects a broadband rough residue $e_c = w_c * \delta$ that spikes V on these channels (A2/A3). P1 ensures norm-preserving tricks cannot suppress the spike. Viyog detects adversarials by thresholding V ; OOD inputs fall below the threshold because A1 applies to all natural inputs.

IV. The Viyog Method

A. Step-by-Step Derivation and Motivation

The first design choice—reading at the first convolutional layer—follows directly from the theory of Section III; the remaining three (restricting to the dormant band, using a total-variation shape statistic, and ID-only threshold calibration) are developed in the subsections that follow:

First convolutional layer. The first convolutional layer is where the affine approximation holds most cleanly [2], [3], the per-channel bound (Eq. (3) in Appendix A) is exact, and—crucially—the adversarial residue has not yet been transformed or pooled. The shape statistic V peaks at the first conv (mean T2 0.965) and decays with depth; the first layer is therefore simultaneously the cheapest and the best layer for the deployed read (Fig. 10).

Given this shape-based score, the routing rule is straightforward.

B. Decision Rule and Downstream Routing

Given a pool of non-ID inputs flagged by a first-stage detector, Viyog applies a second-stage partition:

- $\text{Viyog}(x) > \tau$: ADV — apply adversarial defence, rejection, or security alert.
- $\text{Viyog}(x) \leq \tau$: OOD — apply abstention, fallback, or incremental-learning trigger.

A magnitude threshold allows precision–recall trade-offs. For the deployed Viyog statistic (Sec. IV-C) the OOD-vs-ADV threshold τ is calibrated ID-only (adversarials occupy the high-TV tail), so no held-out ADV data is needed (Alg. 1); the end-to-end gate uses the same statistic (Sec. IX-A).

C. Dormant-Band Shape Statistic (Viyog)

The deployed statistic $V(x) = \text{Viyog}(x)$ uses the dormant band of Eq. (1). Restricting it to active channels is essential: up to 25/64 first-conv channels are permanently silent (e.g. DenseNet-121) and would otherwise force a degenerate score. Across the 37-statistic battery V tops both the ID-vs-ADV and OOD-vs-ADV rankings (Fig. 4).

D. Algorithm and Complexity

Fig. 5 provides a step-by-step visual walkthrough of the full Viyog pipeline with activation-map images at each stage.

Algorithm 1 summarises the deployed Viyog pipeline: calibration computes the per-channel ID profile and dormant band $\mathcal{B}$ once (a single calibration pass), and inference adds only $\mathcal{O}(|\mathcal{B}| h_1 w_1)$ operations and $\mathcal{O}(C)$ stored state (≈ 0.3 KB)—a negligible fraction of the forward pass and four to five orders of magnitude below the megabyte-scale state of feature-distance detectors. We now validate it against eleven baselines across twenty model–dataset configurations.




Fig. 4: First-conv signature ranking by ID-vs-ADV (T2) AUROC on ResNet-50 (37 candidate statistics), coloured by statistic family. Magnitude/norm statistics (orange)—including the raw L_∞ norm (the original Viyog L_∞ -form baseline, annotated)—sit near chance (0.50–0.53). Shape statistics (green)—total-variation (TV) and high-frequency (HF) roughness on the dormant channel band—are best-in-class (≈ 0.98), with Viyog (TV dormant-band) ranked first. Distance/feature (blue), logit (yellow), and distributional (purple) statistics lie in between. This ranking motivates replacing the raw L_∞ -form with the TV dormant-band shape: it is channel shape, not magnitude, that separates adversarial from natural inputs at the first layer.

V. Experimental Setup

A. Datasets and Models

All experiments use two primary ID datasets: CIFAR-100 [36] (224×224 , the main benchmark) and the safety-critical GTSRB [37] (224×224). Per-dataset OOD sets are drawn from SVHN [38], iSUN, LSUN, and Places365; MNIST serves as a supplementary near-OOD probe for DenseNet.

The model panel comprises 20 architectures spanning every major design family: ResNet-18/34/50/101/152 [31], DenseNet-121/161/169/201 [32], ConvNeXtV2 [33], ViT-B/16 [34], Swin-T [35], MobileNetV3-L, MobileNetV4-M, MobileOne-S1, EfficientNet-Lite0, EfficientNetV2-L, EfficientViT-B1, FastViT-SA12, and EdgeNeXt-Small—all fine-tuned from ImageNet pretraining in PyTorch [39] without adversarial augmentation. Training-seed confidence intervals (≤ 3 seeds per architecture) are reported throughout; per-architecture std on T2 is ≤ 0.004 (Fig. 8).

B. Attacks, Baselines, and Metrics

Four L_∞ attacks span the threat model: FGSM [9], BIM [40], PGD [10], and APGD-CE [41] via TorchAttacks [42] at $\varepsilon=8/255$. Eleven pytorch-ood [43] baselines—

How Viyog works — one forward pass, no extra parameters




Fig. 5: How Viyog works, in five steps (one forward pass, no extra parameters). (1) Hook the first-conv activation map $a = f_0(x)$. (2) Rank channels by mean ID activation; the quietest 10% form the dormant band $\mathcal{B}$ (green). (3) Read spatial roughness (total variation) on $\mathcal{B}$: clean channels are smooth (low TV), adversarial channels are jagged (high TV). (4) Score $V(x)$, the mean normalized TV over $\mathcal{B}$: ID and OOD stay low, ADV spikes past the ID-calibrated threshold τ . (5) Route the input—ID passes, OOD abstains, ADV is rejected. Full method and routing rule in Sec. IV (Alg. 1).

Algorithm 1 Viyog: Dormant-band shape statistic for OOD-vs-ADV separation

Require: Model f with first affine layer f_0 ; calibration set $\mathcal{D}$; band fraction $\rho=0.1$; test sample x

- 1: Calibration (once per model):
- 2: $p_c \leftarrow \frac{1}{|\mathcal{D}|} \sum_{x \in \mathcal{D}} |f_0(x)_c|$ $\triangleright$ per-channel ID mean activation
- 3: $\mathcal{A} \leftarrow \{c : p_c > 10^{-4}\}$ $\triangleright$ drop permanently-dead filters
- 4: $\mathcal{B} \leftarrow$ bottom- ρ of $\mathcal{A}$ by p_c $\triangleright$ dormant band
- 5: $\tau \leftarrow$ upper-quantile of s on $\mathcal{D}$ at a target 5% ID-FPR $\triangleright$ ID-only; no ADV labels
- 6: Inference (per sample, single forward pass):
- 7: $a \leftarrow f_0(x)$ $\triangleright$ hook; no extra pass, no gradients
- 8: $s \leftarrow \frac{1}{|\mathcal{B}|} \sum_{c \in \mathcal{B}} \widetilde{\text{TV}}(a_c)$ $\triangleright$ dorm-band roughness, Eq. (1)
- 9: if $s > \tau$ then
- 10: return ADV (reject / defend) $\triangleright$ high-frequency dorm-band residue
- 11: else
- 12: return OOD (abstain / incremental learn)
- 13: end if

Optional 3-way classifier: fit a Linear Discriminant Analysis (LDA) on a small labelled ID/OOD/ADV calibration set over the feature vector $[\text{logit score}, V(x), \dots]$. This adds <0.1 KB of state and returns a three-class ID/OOD/ADV decision (Sec. VI). The unsupervised threshold above is the deployed default.

TABLE II: Experimental coverage—everything evaluated, all under the directionless-AUROC protocol.

Experiment	Coverage	Measures
Detection T1/T2/T3	3 datasets $\times$ 20 archs $\times$ 4 attacks $\times$ 11 detectors	AUROC, FPR@95, AUC
Measured cost	CIFAR-100, H200 (CUDA events)	latency, mem, energy
Layer-depth sweep	6 archs $\times$ 2 attacks $\times$ all depths	T2/T3 depth vs.
Adaptive A0–A3, EOT	20 archs; norm/TV- aware + stochastic	evasion %, AUROC
Strong adaptive A4	6 archs (C100+GTSRB), held- out basis	detector floor
Training-seed CIs	20 archs $\times$ ≤ 3 seeds	T2/T3 std
Non-vision sanity	1D signals, FGSM/PGD	T2 AUROC

where post-hoc tuning is infeasible). We evaluate three deployment tasks: T1 ID-vs-OOD, T2 ID-vs-ADV (the security task), and T3 OOD-vs-ADV (the second-stage routing task). A score is credited only when its separation direction is consistent across all architectures (directionless AUROC [6]); a score whose direction flips model-to-model is marked $\times$ (not deployable).

Experimental coverage. Table II lists every experiment run; all use the same directionless-AUROC protocol. Beyond the main detection sweep we add a layer-depth sweep (Fig. 10), a measured on-silicon cost benchmark against all eleven baselines (Figs. 12, 13), a 20-architecture adaptive ladder with an EOT stochastic-band defence (Sec. VII, Fig. 14), full-20 training-seed CIs (Fig. 8), and a 1D non-vision sanity check (Appendix B).

VI. Results and Discussion

A. Detection Quality Across 20 Architectures

With the evaluation protocol of Sec. V fixed, we first report raw detection quality across all 20 backbones.

MSP, Entropy, MaxLogit, Energy, GEN, ODIN ($T=1000$), MCD (30 passes), Mahalanobis, KNN, OpenMax, KL-Matching—are compared on AUROC, FPR@95TPR (false-positive rate at 95% true-positive rate), and AUC (area under the threshold curve, critical for deployment




Fig. 6: Detection quality across the 20 CIFAR-100 architectures, by task and method family. Viyog leads on both security tasks—T2 0.966 (ID-vs-ADV) and T3 0.824 (OOD-vs-ADV)—while logit detectors win the T1 ID/OOD task (0.85) and feature-distance detectors trail on T3 (0.69–0.71). Tasks are scored directionlessly, not the earlier far-OOD single-direction average. At a strict 5%-FPR point mean ADV recall is 0.345 (median 0.278)—near-zero on ViT-B/EdgeNeXt, a caveat the cascade (Sec. IX-A) addresses.

Under the strict directionless 20-architecture protocol (Table III, Fig. 6), Viyog leads on both the security task T2 (0.966) and the deployment task T3 (0.824) at ≈ 0.3 KB; every logit baseline is near-chance on T2; a three-way supervised head reaches 0.798 balanced accuracy (Fig. 7). The T3 signal is consistent across all 20 architectures, four attacks, and three datasets (Fig. 9); GTSRB confirms the result is architecture- and dataset-stable.

The failure modes are complementary: logit scores collapse on T2 (blind to high-confidence adversarials), while feature-distance detectors (Mahalanobis, KNN) trail on T3, having pulled perturbations back toward the ID manifold at deep layers [44]. Only the first-conv shape statistic separates both (Sec. III-A).

The headline separation is a property of the dormant-band statistic, not a lucky seed; the reproducibility concern is resolved at full-panel scale (Fig. 8).

B. Near-OOD Difficulty and Shape-Depth Optimality

Having established overall separation, we break the T3 result by OOD difficulty and confirm that the first conv is globally optimal for the shape read. ID and OOD produce smooth dormant channels; adversarial perturbations leave high-frequency residue—the shape statistic separates ADV from ID/OOD at $O(C)$ cost (Fig. 16, Appendix A).

Harder near-OOD and texture splits. Because near-distribution or texture-biased OOD may themselves perturb early activations, we break the 17-architecture OOD-vs-ADV (T3) result by difficulty. Deployed Viyog holds far-OOD (0.96) and degrades gracefully but stays above chance on near-OOD (0.74 mean, up to 0.95) and texture-OOD (0.67); near-OOD is indeed the harder regime. The residual gap is complementary (Fig. 11): a high-frequency dorm-band read recovers texture-OOD to 0.92 where the total-variation read is weakest, so pairing Viyog with a second early-layer statistic closes the texture case. The

TABLE III: New directionless re-evaluation across 20 CIFAR-100 architectures (replaces the original far-OOD per-row table). Setup: directionless evaluation across 20 architectures (T1/T2/T3 as defined in Sec. V). Metric: mean AUROC / FPR@95TPR (%). $\times$ = direction-inconsistent (score cannot be deployed); best AUROC per column in bold; shaded row = deployed detector; state memory in last column. Highlights: every logit baseline is direction-inconsistent and near-chance on the security task T2, and the original raw Viyog L_∞ is itself direction-inconsistent here, whereas the deployed Viyog is best on both T2 (0.966) and the deployment task T3 (0.824) at ≈ 0.3 KB. Impact: the deployable signal is carried by first-conv shape, not magnitude or logits.

Method	AUROC / FPR@95TPR (%)			State
	T1 ID/OOD	T2 ID/ADV	T3 OOD/ADV	
Logit baselines (pytorch-ood)				
Energy	0.847 / 59.8	0.414 $\times$ / 98.3	0.789 / 81.6	4B
MSP	0.736 $\times$ / 58.0	0.360 $\times$ / 88.7	0.795 / 81.1	4B
MaxLogit	0.767 $\times$ / 60.0	0.384 $\times$ / 88.8	0.793 / 80.8	4B
Entropy	0.745 $\times$ / 58.1	0.357 $\times$ / 88.9	0.801 / 79.6	4B
GEN	0.745 $\times$ / 59.3	0.340 $\times$ / 88.9	0.809 / 78.8	4B
KL-Match	0.258 $\times$ / 59.2	0.623 $\times$ / 88.6	0.798 / 79.3	4B
First-layer statistics (ours)				
FirstConv- L_∞ (baseline)	0.602 $\times$ / 88.8	0.559 $\times$ / 80.5	0.448 $\times$ / 75.7	8B
Viyog-HF (high-freq. dorm.)	0.759 / 75.8	0.937 / 28.0	0.754 / 61.4	0.3KB
Viyog (deployed)	0.690 / 87.0	0.966 / 13.9	0.824 / 45.5	0.3KB

TABLE IV: Second-stage detector cost (architecture-derived). State is $O(C)$ for the first-layer detectors versus $O(D^2)/O(ND)$ for feature-distance methods. The optional “LDA head” is a tiny linear discriminant fit over the detector scores that turns the binary detector into a 3-class ID/OOD/ADV classifier ($K=3$), adding <0.1 KB and no extra forward pass.

Detector	State	Extra fwd	Order
Energy/MSP (logit)	4 B	0	$O(1)$
FirstConv- L_∞ (baseline)	8 B	0	$O(1)$
Viyog (deployed)	≈ 0.3 KB	0	$O(C)$
+ LDA head (3-class, $K=3$)	<0.1 KB	0	$O(CK)$
Mahalanobis / ViM	4.5–17 MB	1	$O(D^2)$
KNN	20–40 MB	1	$O(ND)$

first conv is globally optimal for V : on the six-architecture depth sweep the shape read peaks at depth 0 (mean T2 0.965, T3 0.918; the slightly higher T3 than the 20-arch panel mean 0.824 reflects the easier six-model subset) and decays with depth, while a raw-norm oracle across all depths still scores below it (Fig. 10)—no adaptive layer selector is needed.

C. Computational Efficiency

Having shown the first-conv read is the most accurate detector, we now show it is also the cheapest.

Figs. 12–13 and Table IV quantify the cost: Viyog uses 0.3 KB and 2.28% of CNN MACs with no extra forward pass, runs 6–347 $\times$ faster than all baselines (H200, CUDA events), and draws 3.8% of inference latency and 12% of energy — the best-separating detector is also the cheapest by orders of magnitude.

Together the AUROC, AUTC, latency, and memory results show Viyog’s first-layer design gives the best ADV–




Fig. 7: Combining a logit detector with Viyog yields three-way ID/OOD/ADV separation—each covers the other’s blind spot. Logit detectors (Energy) catch low-confidence OOD but are blind to high-confidence adversaries; Viyog catches adversaries (dormant-band roughness) but is weak on smooth OOD. (a) Per-class recall (CIFAR-100, 20 archs): Energy alone 0.45 on ADV, Viyog alone 0.29 on OOD; the Energy + Viyog pair raises 3-way balanced accuracy to **0.757** and the full panel to **0.798**. (b) Gain transfers to GTSRB: full panel **0.863**, logit+Viyog 0.768; bootstrap 95% CIs are disjoint (full-panel [0.77,0.83] vs. Energy-only [0.58,0.65]). Viyog adds adversarial detection to any deployed logit OOD detector with no retraining, no extra forward pass, 0.3KB state.




Fig. 8: Training-seed reproducibility across the full 20-architecture panel. Per-architecture mean $\pm$ std (circles), individual seeds (faint crosses), panel mean (vertical line), ± 1 overall std (band). (a) T2 seed std ≤ 0.004 across all 20 backbones—an order of magnitude below the cross-architecture spread (ViT-B the semantic-token outlier). (b) T3 similarly stable (std ≤ 0.06).

OOD separation at the lowest cost. We next stress-test the detector against an adaptive adversary.

VII. Adaptive Attack Evaluation

A rigorous adversarial evaluation must include an adaptive adversary who knows the detection mechanism and specifically targets it.

Threat model. The attacker aims at both misclassification and detector evasion under L_∞ , $\varepsilon=8/255$. Fig. 14 shows the full A0–A3 knowledge ladder; the defender’s hook, band $\mathcal{B}$, and threshold τ are fixed at deployment.

A. Adaptive Attack Formulation

A norm-aware attacker knowing that Viyog monitors first-layer L_∞ norms would optimise perturbations to preserve these norms while inducing misclassification. We formalise this as an augmented PGD objective:

$$\mathcal{L}_{\text{adap}} = \mathcal{L}_{\text{CE}}(f(x+\delta), y) + \lambda (\|f_0(x+\delta)\|_\infty - \|f_0(x)\|_\infty)^2, \quad (2)$$

where $\lambda > 0$ controls the penalty strength; we sweep L_∞/L_2 /combined modes ($\varepsilon=8/255$, $\alpha=2/255$, 10 PGD steps) and confirm all ten standard attacks suppress first-layer norms without special treatment.

B. Results and Security Implications

The ladder of Fig. 14 measures what each rung of attacker knowledge costs—from black-box PGD (A0) through a both-aware adaptive adversary (A3) to a stochastic-band defence (EOT).

With the penalty active, $\|f_0(x_{\text{adap}})\|_\infty \approx \|f_0(x)\|_\infty$, keeping the raw score in-distribution, yet the dormant-band shape read is unmoved. The A0–A3 ladder (Fig. 14) shows no single statistic is load-bearing: norm-preserving A1 gives 0% evasion (P1), and only the both-aware A3 suppresses every signature—at a measurable 8% attack-success cost.

The result holds on a second dataset. On GTSRB (ResNet-50, ConvNeXtV2, ViT-B) at max penalty λ , the both-aware attacker keeps $\geq 80\%$ success yet the combined max(dorm, HF) detector never drops below AUROC 0.775: single-channel evasion is always backstopped by the complement (ResNet-50 dorm 0.83 $\rightarrow$ 0.51 but HF stays 0.72; ConvNeXtV2 dorm-suppression costs 80% of success). The attacker thus either keeps the detector ≥ 0.775 or surrenders most of its success—on both datasets.

Evasion requires a unique and costly modification. All these standard (A0) attacks [9]–[12], [41] naturally suppress first-layer norms and are detected; only the bespoke penalty of Eq. (2) defeats the norm, a far higher barrier than off-the-shelf methods [42]. This is not an empirical quirk: by the first-layer mechanism (Sec. III-A), any gradient-based optimisation suppresses the high-influence channels through the Jacobian, so the attacker must explicitly counteract it with an added loss term.

L_∞ evasion yields a safer failure mode. Evading the global L_∞ score—the only signature a norm-preserving attack defeats (above)—lifts the sample’s norm into the ID or OOD regime, so the system classifies it as ID/OOD




Fig. 9: The OOD-vs-ADV (T3) signal is consistent across architectures, attacks, and datasets. (a) Per-architecture T3 AUROC on all 20 CIFAR-100 backbones (mean 0.82, range 0.65–0.95, every backbone above chance). (b) T3 across four attacks (FGSM, BIM, PGD, APGD-CE) and three datasets, weakest case ≥ 0.77 . (c) Viyog matches the strongest logit detector (GEN) within ± 0.02 at 0.3KB and no extra forward pass, yet stays complementary to it (Fig. 7).




Fig. 10: The first conv is the optimal layer for the deployed shape read—adaptive layer selection is unnecessary. ID-vs-ADV (T2) AUROC of the dormant-band shape statistic V and of the raw L_∞ statistic at every depth, averaged over six architectures (ResNet-50, DenseNet-121, MobileNetV3, ConvNeXtV2, Swin-T, ViT-B) and two attacks (PGD, APGD). The shape read (green) peaks at the first conv (mean 0.965) and decays with depth, whereas the raw norm (red) is weakest there (0.52) and peaks only deep (0.80, depth 6). The first-conv shape read beats the best raw-norm layer at any depth—the cheapest layer is also the best.

and abstains or seeks review rather than silently acting on a corrupted prediction—strictly safer than an undetected adversarial.

C. The Robustness–Evasion Trade-Off

The norm-preservation and misclassification objectives compete: budget spent preserving the convolutional norm is budget unavailable for misclassification, so matching standard-PGD success under the constraint requires a larger ϵ or more iterations—the measurable 8% success cost of the both-aware attack above. This trade-off is fundamental to the shared Jacobian structure of Eq. (3) and cannot be circumvented by gradient masking. A stronger held-out-basis (A4) attacker that re-optimises its perturbation basis directly against the deployed statistic does not change this conclusion (see Appendix A, Fig. 17): the same trade-off holds, with $\sim 15\%$ of attack success surrendered to floor both signatures above chance. With adaptive robustness established, the next section shows the norm-suppression phenomenon generalises beyond CNNs.

VIII. Generalisation Beyond CNNs

The theoretical argument of Section III requires only a near-linear first projection and gradient-based norm suppression; it is not specific to convolutional architectures. We verify empirically on three architecturally distinct models fine-tuned on CIFAR-100 at 224×224 resolution from ImageNet-21k pretrained weights.

ConvNeXtV2-Base [33], ViT-B/16 [34], and Swin-T [35] span depthwise-convolutional, pure self-attention, and hi-


Complementary reads cover every OOD difficulty (mean $\pm$ std, 17 archs)

Fig. 11: Viyog is a complementary panel, not a single statistic: paired early-layer reads separate ADV from every OOD difficulty. OOD-vs-ADV (T3) AUROC on the 17-architecture CIFAR-100 panel (mean $\pm$ std), split by OOD difficulty: far (a different dataset), near (a semantically similar shift), and texture (texture-biased). The deployed dormant-band total-variation read (Viyog-D, green) is strongest on far/near-OOD but weakest on texture; its high-frequency complement (HF, blue) is the mirror image. Taking the per-architecture better of the two (Panel, amber) clears 0.76–0.97 across all three difficulties, recovering texture-OOD from 0.67 to 0.92, while the raw L_∞ magnitude norm (grey) stays near chance (0.5) throughout—the separable signal is first-conv shape, not magnitude.

erarchical shifted-window designs; in every case the first layer is an affine strided-convolution or patch-embedding map (ConvNeXtV2 4×4 conv, ViT-B 16×16 and Swin-T 4×4 patch embeddings) satisfying Eq. (3) exactly at the pre-activation stage (the deployed statistic is read post-activation and carries the bounded GELU/ReLU remainder of Appendix A).

On all three, the deployed dormant-band roughness $V(x)$ separates ADV from ID/OOD as in the 32 × 32 CNN results, and the separation is seed-invariant (per-architecture training-seed CIs, Sec. IX-B) and stable across datasets under evaluation resampling (Fig. 8), addressing the reliability concern of Szyk et al. [6].

The generalisation is therefore structural: every model begins with a learned near-linear first projection, so the signal survives even the implicit robustness of large-scale pretraining and is not an artefact of fragile undefended models. Having established the detector’s accuracy, cost, adaptive robustness, and cross-architecture reach, we now turn it into an operational three-way classifier.

IX. Downstream Utility: Three-Way Classification

Viyog turns a binary anomaly detector into a three-way ID/OOD/ADV classifier, an operational distinction that diagnostic metrics (AUROC, FPR) do not capture: without it a system must apply one response to all non-ID inputs—wasting adversarial purification [18]–[20] or feature-denoising [16] on benign OOD, or leaving attacks

TABLE V: End-to-end two-stage recall at a matched 5% ID false-positive rate (mean over 17 archs × 10 seeds). A two-signal gate that adds the first-conv deviation closes the adversarial-recall gap with zero mis-escalation.

Stage-1 gate	ADV recall	OOD recall	ID→ADV
Energy only	0.156	0.343	0.002
Energy $\oplus$ dorm-z (2-signal)	0.835	0.412	0.000

undetected. Beyond these binary routing cases, in class-incremental learning [45], [46] OOD inputs should feed back as new classes while ADV inputs must not contaminate training.

A. End-to-End Cascade and False-Positive Propagation

We evaluate the full two-stage pipeline (non-ID gate → Viyog) over 17 architectures × 10 seeds at a matched 5% ID-FPR (Table V, Fig. 15).

A logit (Energy) gate is itself blind to adversarials—they are confidently wrong—so end-to-end ADV recall is only 0.156. Adding the near-free first-conv deviation to the gate ($\max(z_{\text{energy}}, |z_{\text{dorm}}|)$, same 5% ID-FPR) lifts it to **0.835** (mean over 16/17 architectures; the weakest are ViT-B 0.21—its patch-embedding has no dormant low-activation band—fastvit 0.62 and ResNet-50 0.73, while OOD recall holds at ≈ 0.41) with ID→ADV mis-escalation → 0: one first-conv statistic both flags non-ID ($|z|$) and routes OOD-vs-ADV (sign). For false-positive propagation, Viyog escalates a wrongly-flagged ID input to the ADV class only **5.2%** of the time (median 0; exactly 0 on 14/17 architectures) versus 51–82% for logit/ L_∞ —mitigating rather than amplifying upstream errors.

B. Limitations and Future Work

While Viyog addresses the core OOD-vs-ADV separation problem, six boundary conditions remain. Three are resolved in this revision (multi-seed reproducibility, non-vision generality, and adaptive robustness); two are open by design scope (physical attacks attenuated by the camera ISP, and on-board edge measurements); and one is a benign artefact (DenseNet-121 dead-channel degeneracy on MNIST, fixed by active-channel selection). Table VI summarises each item with its current status.

X. Conclusion

We proposed Viyog, a lightweight gradient-free method for separating OOD from ADV inputs via first-layer dormant-band roughness statistics. The deployed detector is the dormant-band roughness statistic $V(x)$, grounded in Proposition 1: it inflates under the exact broadband residue any gradient attack injects while remaining invariant to activation magnitude, separating ADV from ID/OOD at AUROC **0.966** (T2), confirmed stable across 20 architectures and 3 training seeds (Table III, Fig. 8). The phenomenon persists in ConvNeXtV2, ViT-B/16, and Swin-T at 224 × 224; a principled adaptive analysis shows defeating it requires an explicit norm-preservation loss absent from standard attacks, and even then only the global




Fig. 12: Measured second-stage detection cost versus every pytorch-ood baseline (CIFAR-100; mean over ResNet-50, MobileNetV3 and DenseNet-121, wall-clock with CUDA sync on one H200). (a) Per-image detection latency (log scale): Viyog reads only the first conv plus an $O(C)$ reduction (0.05 ms/img), so it is $\sim 6\times$ faster than the logit detectors MSP/Energy, 9–21 $\times$ faster than ViM/Mahalanobis, 140 $\times$ faster than KNN, and 347 $\times$ faster than MC-Dropout (30 forward passes). (b) Latency–memory Pareto (log–log): Viyog sits alone in the lower-left corner, storing 0.23 KB against the 9–28 MB covariance/feature-bank of the distance detectors (10^4 – $10^5\times$ smaller). Impact: the detector with the best ADV–OOD separation (Table III) is also the cheapest to run and to store, by orders of magnitude — and the time advantage is now measured, not inferred from complexity order.




Fig. 13: Measured on-silicon detector cost and edge projection (closes the on-device request). (a) Per-architecture cost of the deployed $V(x)$ detector as a fraction of full inference: the exact MAC fraction (mean 2.28%) and the measured batch-1 latency on an H200 (CUDA events; mean 3.8%) agree, confirming a single-digit-% second-stage tax on real silicon across CNNs, a ViT, and a ConvNeXt. (b) Compute-bound roofline projection of the detector’s per-image energy on real edge accelerators (Jetson Orin Nano, RPi5 + Hailo-8L): 19–37 μ J versus a 0.6–1.2 mJ full inference. Energy scales with active-compute time, so the energy fraction follows the latency fraction; absolute roofline figures are optimistic lower bounds, but the fraction is device-independent.

L_∞ score is evaded—the complementary dormant-band shape statistic still flags the attack (0% evasion)—while L_∞ evasion itself yields the safer outcome of abstention.

Complementary to logit scores (which are blind to adversarials), Viyog in a lightweight supervised head gives the first deployable three-way ID/OOD/ADV separation—full-panel balanced accuracy 0.798–0.863 across 20 architectures and two datasets, the logit+Viyog pair alone 0.757/0.768 (CIFAR-100/GTSRB), at ≈ 0.3 KB of added state. Viyog adds no parameters, requires no retraining, and runs in a single forward pass, making it immediately deployable as a plug-in second stage for safety-critical embedded AI.

TABLE VI: Limitations and status. Camera-ready items are committed; addressed items are validated in the cited figures/sections.

Limitation	Detail	Status
(i) Seeds	T2 std ≤ 0.004 across all 20 archs	resolved (Fig. 8)
(ii) MNIST	Low-complexity images; affects path only	benign; T2/T3 unaffected
(iii) Physical	camera image-signal processor (ISP) low-pass filter attenuates broadband residue	open; digital only
(iv) Non-vision	1D sanity check: AU-ROC 0.96 (3 seeds)	resolved (App. B)
(v) Adaptive	20-arch sweep; A3 costs 8% success	resolved (Fig. 14)
(vi) On-silicon	H200 measured roofline; no Jetson/Pi board	measured + measured (Fig. 13)

Appendix A

General-Affine Suppression Bound, GELU Error, and Proof of Proposition 1

The first-layer perturbation bound (Eq. (3)) uses only that the first layer is affine. For any first-layer Jacobian W (dense layer, strided convolution, or ViT/Swin patch embedding) and perturbation $\|\delta\|_\infty \leq \varepsilon$,

$$\|W\delta\|_\infty \leq \|W\|_\infty \varepsilon, \quad \|W\|_\infty = \max_i \sum_j |W_{ij}|, \quad (3)$$

the induced ∞ -operator norm (maximum absolute row sum). Convolution is the special case in which $W = M_m$ is Toeplitz and the maximum row sum equals $\|W_m\|_1$, recovering Eq. (3); the Toeplitz structure is thus sufficient but not necessary, so the mechanism extends to patch-embedding first layers that have no shift-invariance.

If the first activation is a smooth g (e.g. GELU), the linearisation error in each pre-activation coordinate is bounded by the Taylor remainder




Fig. 14: Defence-in-depth against an escalating adaptive adversary (mean over 20 architectures). Attacker knowledge grows left-to-right: A0 (PGD), A1 norm-preserving, A2 TV-aware (white-box), A3 both-aware, EOT stochastic defence. (a) A1 achieves 0% evasion (P1); A2 erodes V to 0.64 while the complement holds at 0.74; EOT restores V to 0.82; A3 suppresses both only at an 8% attack-success cost. (b) EOT scatter: 18/20 architectures above the diagonal confirm stochastic band-randomisation as the reliable system-level backstop.




Fig. 15: End-to-end two-stage cascade over all 17 architectures at a matched 5% ID-FPR. (a) Energy-only gate: mean ADV recall 0.16; two-signal gate ($\max(z_{\text{energy}}, |z_{\text{dorm}}|)$, where z_X denotes the normalised score for statistic X): mean 0.84. (b) ID→ADV mis-escalation stays ≤ 0.001 .

$\frac{1}{2} \sup_z |g''(z)| (\|W\|_{\infty} \epsilon)^2$. For $\text{GELU}(z) = z \Phi(z)$, $g''(z) = \varphi(z)(2 - z^2)$ has interior extrema at $z = 0$ (value $2\varphi(0)$) and $z = \pm 2$ (value $-2\varphi(2) \approx -0.108$); since $0.798 > 0.108$ the supremum is at $z = 0$,

$$\sup_z |\text{GELU}''(z)| = 2\varphi(0) = \sqrt{2/\pi} \approx 0.798, \quad (4)$$

so the per-coordinate error is at most $0.399(\|W\|_{\infty} \epsilon)^2$. This is a worst-case bound: the remainder-to-linear ratio is $0.399\|W\|_{\infty} \epsilon$, which for the measured first-layer weights is $O(1)$ – $O(10)$ (e.g. ≈ 6.8 for ResNet-50 at $\epsilon=8/255$). The bound therefore certifies the affine model only to leading order rather than as an exact identity; the realised perturbation reaches only a fraction of this operator-norm bound (Sec. III), but we do not claim the post-activation linearisation is exact.

Full Proof of Proposition 1: The affine residue $e_c = w_c * \delta$ established above is the object on which the roughness




Fig. 16: The dormant-band signature (supplementary). Per-image dormant-band TV profile for ID (green), OOD (blue), and adversarial (red) inputs on ResNet-50 (CIFAR-100, five OOD sets, PGD attack). ID and OOD share a low, smooth profile; adversarials inject high-frequency structure, separating at AUROC 0.91.

statistic acts; we now prove its three properties (P1–P3) in turn.

(P1) Scale-invariance. TV and $\overline{|\cdot|}$ are absolutely 1-homogeneous, so $\text{TV}(\lambda_c a_c) = \frac{\lambda_c \text{TV}(a_c)}{\lambda_c \nu_c + \epsilon}$, equal to $\text{TV}(a_c)$ at $\epsilon = 0$; first-order expansion in $\epsilon/(\lambda_c \nu_c)$ gives the stated $O(\epsilon/\nu_c)$ residual. V averages such terms, inheriting the invariance. An attacker holding any magnitude norm fixed therefore cannot move V ; to lower V it must smooth the injected residue e_c , which (A2) conflicts with flipping the label.

(P2) Two-sided bound. Affineness gives $e_c = w_c * \delta$ exactly. Upper bound: TV subadditivity $\text{TV}(a'_c) \leq \text{TV}(a_c) + \text{TV}(e_c)$; divide by $\nu_c + \epsilon$, use $\text{TV}(e_c) = \rho_c r_c \nu_c$ (A2, A3) and $\text{TV}(a_c) \leq \beta_c \nu_c$ (A1). Lower bound: reverse triangle $\text{TV}(a'_c) \geq \text{TV}(e_c) - \text{TV}(a_c)$; the denominator satisfies $|a'_c| \leq \nu_c(1 + r_c)$; substitute A1–A3 and cancel ν_c . The Young’s convolution envelope $\text{TV}(e_c) \leq \|w_c\|_1 \|\nabla \delta\|_1$ is




Fig. 17: Strong adaptive frontier—held-out basis (6 archs). An A4 attacker that re-optimises its perturbation basis directly. Dorm/HF mean 0.86/0.79 (min 0.54/0.53); only the both-aware objective suppresses both to ≈ 0.62 at a 15% attack-success cost, confirming the A3 trade-off holds at stronger attack strength.




Fig. 18: Dormant-band shape transfers to 1D signals. ID-vs-ADV AUROC of $V(x)$ (green) vs. raw L_∞ (red) at the first Conv1d layer for FGSM/PGD at two budgets ($n=1,500$, 3 seeds). Shape holds at 0.95–1.00 while the raw norm stays near chance.

$\sim 2\times$ loose empirically, so the norm bound is not tight—the separation comes from the spectral-shape contrast $\rho_c - \beta_c$. For post-ReLU/GELU layers (1-Lipschitz, monotone), $|\Delta\phi| \leq |\Delta u|$ pointwise, so the TV inequalities carry through unchanged.

(P3) Conditional separation. The per-channel margin is $\frac{\rho_{\min} r_{\min}}{1+r_c} - 2\beta_c$, positive under the stated condition; averaging over $\mathcal{B}$ yields $V(x') > V(x)$. The margin grows with r_c (the quietness ratio), formalising why dormant channels maximise separation and why dead-channel exclusion (κ) is necessary. Why OOD stays low: OOD images are natural (A1 applies) and P1 ensures magnitude changes cannot raise V —exactly the failure mode of the raw- $\|a\|_\infty$ statistic that the shape ratio avoids.

Appendix B

Non-Vision Sanity Check (1D Signals)

Because the suppression bound of Appendix A assumes only an affine first layer, the mechanism should transfer to any modality with one; we verify this on 1D signals. A 1D-CNN trained on band-limited multi-sinusoid signals detects adversarials at mean T2 AUROC **0.96** (vs. 0.71 for

the raw norm, 3 seeds; Fig. 18), confirming the broadband-roughness mechanism is modality-agnostic.

Acknowledgment

This research was supported by Global – Learning & Academic research institution for Master’s · PhD students, and Postdocs (G-LAMP) Program of the National Research Foundation of Korea (NRF) grant funded by the Ministry of Education (No. RS-2025-25442252), the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (RS-2025-00561169), and the Institute of Information & communications Technology Planning & Evaluation (IITP) under the Leading Generative AI Human Resources Development (IITP-2026-RS-2026-25546026) grant funded by the Korea government (MSIT).

This research was supported by the Ministry of Science and ICT (MSIT), Korea, under the National Program in Medical AI Semiconductor (No. 2024-0-0097), supervised by the Institute of Information & Communications Technology Planning & Evaluation (IITP) in 2026.

This work was supported by an Institute of Information & Communications Technology Planning & Evaluation (IITP) grant funded by the Korean government (MSIT) (No. RS-2025-02304331, Digital Columbus Project).

This work was partially supported by funding from National Science Foundation grants ACED 2436016.

This research was supported by the ANCHOR Program through the Gangwon ANCHOR Center, funded by the Ministry of Education (MOE) and Gangwon State (G.S.), Republic of Korea (No. 2026-RISE-10-006).

References

- [1] J. Perez-Cerrolaza, J. Abella, M. Borg, C. Donzella, J. Cerquides, F. J. Cazorla, C. Englund, M. Tauber, G. Nikolakopoulos, and J. L. Flores, “Artificial intelligence for safety-critical systems in industrial and transportation domains: A survey,” *ACM Comput. Surv.*, vol. 56, no. 7, Apr. 2024. [Online]. Available: <https://doi.org/10.1145/3626314>
- [2] J. Suresh, N. Nayak, and S. Kalyani, “First line of defense: a robust first layer mitigates adversarial attacks,” in *Proceedings of the Thirty-Ninth AAAI Conference on Artificial Intelligence and Thirty-Seventh Conference on Innovative Applications of Artificial Intelligence and Fifteenth Symposium on Educational Advances in Artificial Intelligence*, ser. AAAI’25/IAAI’25/EAAI’25. AAAI Press, 2025. [Online]. Available: <https://doi.org/10.1609/aaai.v39i7.32771>
- [3] C. Bakiskan, M. Cekic, and U. Madhow, “Early layers are more important for adversarial robustness,” in *ICLR 2022 workshop on new frontiers in adversarial machine learning*, 2022.
- [4] J. Yang, K. Zhou, Y. Li, and Z. Liu, “Generalized out-of-distribution detection: A survey,” *International Journal of Computer Vision*, vol. 132, no. 12, pp. 5635–5662, 2024.
- [5] Z. Hong, Y. Yue, Y. Chen, L. Cong, H. Lin, Y. Luo, M. H. Wang, W. Wang, J. Xu, X. Yang et al., “Out-of-distribution detection in medical image analysis: A survey,” *arXiv preprint arXiv:2404.18279*, 2024.
- [6] K. Szyc, T. Walkowiak, and H. Maciejewski, “Why out-of-distribution detection experiments are not reliable — subtle experimental details muddle the OOD detector rankings,” in *Uncertainty in Artificial Intelligence*, ser. *Proceedings of Machine Learning Research*. PMLR, 2023.
- [7] J. Zhang et al., “Splitting & integrating: Out-of-distribution detection via adversarial gradient attribution,” in *Forty-Second International Conference on Machine Learning*, 2024.

- [8] C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus, "Intriguing properties of neural networks," arXiv preprint arXiv:1312.6199, 2013.
- [9] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," arXiv preprint arXiv:1412.6572, 2014.
- [10] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, "Towards deep learning models resistant to adversarial attacks," arXiv preprint arXiv:1706.06083, 2017.
- [11] N. Carlini and D. Wagner, "Towards evaluating the robustness of neural networks," in 2017 IEEE Symposium on Security and Privacy (SP). IEEE, 2017, pp. 39–57.
- [12] S.-M. Moosavi-Dezfooli, A. Fawzi, and P. Frossard, "Deepfool: a simple and accurate method to fool deep neural networks," in Proceedings of the IEEE conference on computer vision and pattern recognition, 2016, pp. 2574–2582.
- [13] D. Tsipras, S. Santurkar, L. Engstrom, A. Turner, and A. Madry, "Robustness may be at odds with accuracy," arXiv preprint arXiv:1805.12152, 2018.
- [14] H. Zhang, Y. Yu, J. Jiao, E. Xing, L. El Ghaoui, and M. Jordan, "Theoretically principled trade-off between robustness and accuracy," in International conference on machine learning. PMLR, 2019, pp. 7472–7482.
- [15] V. Ziyadinov and M. Tereshonok, "Low-pass image filtering to achieve adversarial robustness," Sensors, vol. 23, no. 22, p. 9032, 2023.
- [16] C. Xie, Y. Wu, L. v. d. Maaten, A. L. Yuille, and K. He, "Feature denoising for improving adversarial robustness," in Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, 2019, pp. 501–509.
- [17] P. Samangouei, M. Kabkab, and R. Chellappa, "Defense-gan: Protecting classifiers against adversarial attacks using generative models," arXiv preprint arXiv:1805.06605, 2018.
- [18] W. Nie, B. Guo, Y. Huang, C. Xiao, A. Vahdat, and A. Anandkumar, "Diffusion models for adversarial purification," arXiv preprint arXiv:2205.07460, 2022.
- [19] J. Yoon, S. J. Hwang, and J. Lee, "Adversarial purification with score-based generative models," 2021. [Online]. Available: <https://arxiv.org/abs/2106.06041>
- [20] H. P. Silva, L. Seidenari, and A. Del Bimbo, "Diffdefense: Defending against adversarial attacks via diffusion models," in International Conference on Image Analysis and Processing. Springer, 2023, pp. 430–442.
- [21] G. S. Dhillon, K. Azizzadenesheli, Z. C. Lipton, J. Bernstein, J. Kossai, A. Khanna, and A. Anandkumar, "Stochastic activation pruning for robust adversarial defense," 2018. [Online]. Available: <https://arxiv.org/abs/1803.01442>
- [22] R. Chen, H. Jin, H. Zheng, J. Chen, and Z. Liu, "Fight perturbations with perturbations: Defending adversarial attacks via neuron influence," IEEE Transactions on Dependable and Secure Computing, vol. 22, no. 2, pp. 1582–1595, 2025.
- [23] X. Li and F. Li, "Adversarial examples detection in deep networks with convolutional filter statistics," in Proceedings of the IEEE international conference on computer vision, 2017, pp. 5764–5772.
- [24] H. Wang, X. Wu, Z. Huang, and E. P. Xing, "High-frequency component helps explain the generalization of convolutional neural networks," in Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, 2020, pp. 8684–8694.
- [25] Z. Wang, Y. Yang, A. Shrivastava, V. Rawal, and Z. Ding, "Towards frequency-based explanation for robust cnn," arXiv preprint arXiv:2005.03141, 2020.
- [26] Y. Feng, D. Jun, X. Ng, and A. Easwaran, "A unified perspective on adversarial and out-of-distribution detection in the open world," in Proceedings of the AAAI Conference on Artificial Intelligence, Workshop on Engineering Dependable and Secure Machine Learning Systems, 2022.
- [27] T. Walkowiak, K. Szyk, and H. Maciejewski, "Combining outlier-ness scores and feature extraction techniques for improvement of ood and adversarial attacks detection in dnns," in International Conference on Computational Science. Springer, 2023, pp. 578–592.
- [28] A. L. Pellicer, K. Giatgong, Y. Li, N. Suri, and P. Angelov, "Unicad: A unified approach for attack detection, noise reduction and novel class identification," in 2024 International Joint Conference on Neural Networks (IJCNN). IEEE, 2024, pp. 1–8.
- [29] W. Liu, X. Wang, J. Owens, and Y. Li, "Energy-based out-of-distribution detection," in Advances in Neural Information Processing Systems, H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, Eds., vol. 33. Curran Associates, Inc., 2020, pp. 21464–21475. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2020/file/f5496252609c43eb8a3d147ab9b9c006-Paper.pdf
- [30] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," arXiv preprint arXiv:1610.02136, 2016.
- [31] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in Proceedings of the IEEE conference on computer vision and pattern recognition, 2016, pp. 770–778.
- [32] G. Huang, Z. Liu, L. Van Der Maaten, and K. Q. Weinberger, "Densely connected convolutional networks," in Proceedings of the IEEE conference on computer vision and pattern recognition, 2017, pp. 4700–4708.
- [33] S. Woo, S. Debnath, R. Hu, X. Chen, Z. Liu, I. S. Kweon, and S. Xie, "ConvNeXt V2: Co-designing and scaling ConvNets with masked autoencoders," in Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2023, pp. 16133–16142.
- [34] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An image is worth 16x16 words: Transformers for image recognition at scale," International Conference on Learning Representations, 2021.
- [35] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, "Swin transformer: Hierarchical vision transformer using shifted windows," in Proceedings of the IEEE/CVF International Conference on Computer Vision, 2021, pp. 10012–10022.
- [36] A. Krizhevsky, "Learning multiple layers of features from tiny images," University of Toronto, Tech. Rep., 2009.
- [37] J. Stallkamp, M. Schlipsing, J. Salmen, and C. Igel, "Man vs. computer: Benchmarking machine learning algorithms for traffic sign recognition," Neural Networks, vol. 32, pp. 323–332, 2012.
- [38] Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu, A. Y. Ng et al., "Reading digits in natural images with unsupervised feature learning," in NIPS workshop on deep learning and unsupervised feature learning, vol. 2011, no. 2. Granada, 2011, p. 4.
- [39] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga et al., "Pytorch: An imperative style, high-performance deep learning library," Advances in neural information processing systems, vol. 32, 2019.
- [40] A. Kurakin, I. Goodfellow, and S. Bengio, "Adversarial machine learning at scale," arXiv preprint arXiv:1611.01236, 2016.
- [41] F. Croce and M. Hein, "Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks," in International conference on machine learning. PMLR, 2020, pp. 2206–2216.
- [42] H. Kim, "Torchattacks: A pytorch repository for adversarial attacks," arXiv preprint arXiv:2010.01950, 2020.
- [43] K. Kirchheim, M. Filax, and F. Ortmeier, "Pytorch-ood: A library for out-of-distribution detection based on pytorch," in Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2022, pp. 4351–4360.
- [44] K. Lee, K. Lee, H. Lee, and J. Shin, "A simple unified framework for detecting out-of-distribution samples and adversarial attacks," Advances in neural information processing systems, vol. 31, 2018.
- [45] S.-A. Rebuffi, A. Kolesnikov, G. Sperl, and C. H. Lampert, "icarl: Incremental classifier and representation learning," in Proceedings of the IEEE conference on Computer Vision and Pattern Recognition, 2017, pp. 2001–2010.
- [46] X. Cao, H. Lu, X. Liu, and M.-M. Cheng, "Class incremental learning for image classification with out-of-distribution task identification," IEEE Transactions on Multimedia, 2025.